```
!pip install -U transformers
```

```
Requirement already satisfied: transformers in /usr/local/lib/python3.12/dist
Collecting transformers
  Downloading transformers-5.4.0-py3-none-any.whl.metadata (32 kB)
Requirement already satisfied: huggingface-hub<2.0,>=1.5.0 in /usr/local/lib/
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.12/dist-
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/d
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.12/dist-
Requirement already satisfied: regex>=2025.10.22 in /usr/local/lib/python3.12
Requirement already satisfied: tokenizers<=0.23.0,>=0.22.0 in /usr/local/lib/
Requirement already satisfied: typer in /usr/local/lib/python3.12/dist-packag
Requirement already satisfied: safetensors>=0.4.3 in /usr/local/lib/python3.1
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.12/dist-p
Requirement already satisfied: filelock>=3.10.0 in /usr/local/lib/python3.12/
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.12/
Requirement already satisfied: hf-xet<2.0.0,>=1.4.2 in /usr/local/lib/python3
Requirement already satisfied: httpx<1,>=0.23.0 in /usr/local/lib/python3.12/
Requirement already satisfied: typing-extensions>=4.1.0 in /usr/local/lib/pyt
Requirement already satisfied: click>=8.2.1 in /usr/local/lib/python3.12/dist
Requirement already satisfied: shellingham>=1.3.0 in /usr/local/lib/python3.1
Requirement already satisfied: rich>=12.3.0 in /usr/local/lib/python3.12/dist
Requirement already satisfied: annotated-doc>=0.0.2 in /usr/local/lib/python3
Requirement already satisfied: anyio in /usr/local/lib/python3.12/dist-packag
Requirement already satisfied: certifi in /usr/local/lib/python3.12/dist-pack
Requirement already satisfied: httpcore==1.* in /usr/local/lib/python3.12/dis
Requirement already satisfied: idna in /usr/local/lib/python3.12/dist-package
Requirement already satisfied: h11>=0.16 in /usr/local/lib/python3.12/dist-pa
Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/pyth
Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.12/dist-p
Downloading transformers-5.4.0-py3-none-any.whl (10.1 MB)
   ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 10.1/10.1 MB 87.7 MB/s eta 0:00:0
Installing collected packages: transformers
  Attempting uninstall: transformers
    Found existing installation: transformers 5.0.0
    Uninstalling transformers-5.0.0:
      Successfully uninstalled transformers-5.0.0
Successfully installed transformers-5.4.0
```

Local Inference on GPU

Model page: https://huggingface.co/Salesforce/blip-vqa-base

⚠️ If the generated code snippets do not work, please open an issue on either the [model repo](#) and/or on [huggingface.js](#) 🙏

```
from transformers import BlipProcessor, BlipForQuestionAnswering
from PIL import Image
import requests

# Load processor and model
processor = BlipProcessor.from_pretrained("Salesforce/blip-vqa-base")
model = BlipForQuestionAnswering.from_pretrained("Salesforce/blip-vqa-ba
```

```

# Load an image
img_url = "https://storage.googleapis.com/sfr-vision-language-research/B
image = Image.open(requests.get(img_url, stream=True).raw).convert("RGB"

# Ask a question
question = "What is in the image?"
inputs = processor(image, question, return_tensors="pt")

# Generate answer
out = model.generate(**inputs)
print(processor.decode(out[0], skip_special_tokens=True))
```

```
/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: Us
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings ta
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access p
  warnings.warn(
Warning: You are sending unauthenticated requests to the HF Hub. Please set a
WARNING:huggingface_hub.utils._http:warning: You are sending unauthenticated
```

```
preprocessor_config.json: 100%  445/445 [00:00<00:00, 8.91kB/s]

config.json:  4.56k/? [00:00<00:00, 69.3kB/s]

tokenizer_config.json: 100%  592/592 [00:00<00:00, 13.5kB/s]

vocab.txt:  232k/? [00:00<00:00, 2.26MB/s]

tokenizer.json:  711k/? [00:00<00:00, 9.95MB/s]

special_tokens_map.json: 100%  125/125 [00:00<00:00, 3.05kB/s]

model.safetensors: 100%  1.54G/1.54G [00:26<00:00, 75.9MB/s]

Loading weights: 100%  788/788 [00:00<00:00, 3791.23it/s]
```

```
BlipForQuestionAnswering LOAD REPORT from: Salesforce/blip-vqa-base
Key                                       | Status     |  |
------------------------------------------+------------+--+-
text_encoder.embeddings.position_ids      | UNEXPECTED |  |
text_decoder.bert.embeddings.position_ids | UNEXPECTED |  |

Notes:
- UNEXPECTED:   can be ignored when loading from different task/architecture;
/usr/local/lib/python3.12/dist-packages/transformers/generation/utils.py:1554
  warnings.warn(
woman and dog
```

```
# Load model directly
from transformers import AutoProcessor, AutoModelForVisualQuestionAnswer

processor = AutoProcessor.from_pretrained("Salesforce/blip-vqa-base")
model = AutoModelForVisualQuestionAnswering.from_pretrained("Salesforce/
```

```
Loading weights: 100%                                    788/788 [00:00<00:00, 4046.00it/s]
BlipForQuestionAnswering LOAD REPORT from: Salesforce/blip-vqa-base
Key                                        | Status     |  |
-------------------------------------------+------------+--+-
text_encoder.embeddings.position_ids       | UNEXPECTED |  |
text_decoder.bert.embeddings.position_ids  | UNEXPECTED |  |

Notes:
- UNEXPECTED:   can be ignored when loading from different task/architecture;
```

```
!pip install -q --upgrade transformers
```

```
!pip install transformers datasets pillow
```

```
Requirement already satisfied: transformers in /usr/local/lib/python3.12/dist
Requirement already satisfied: datasets in /usr/local/lib/python3.12/dist-pac
Requirement already satisfied: pillow in /usr/local/lib/python3.12/dist-packa
Requirement already satisfied: huggingface-hub<2.0,>=1.5.0 in /usr/local/lib/
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.12/dist-
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/d
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.12/dist-
Requirement already satisfied: regex>=2025.10.22 in /usr/local/lib/python3.12
Requirement already satisfied: tokenizers<=0.23.0,>=0.22.0 in /usr/local/lib/
Requirement already satisfied: typer in /usr/local/lib/python3.12/dist-packag
Requirement already satisfied: safetensors>=0.4.3 in /usr/local/lib/python3.1
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.12/dist-p
Requirement already satisfied: filelock in /usr/local/lib/python3.12/dist-pac
Requirement already satisfied: pyarrow>=15.0.0 in /usr/local/lib/python3.12/d
Requirement already satisfied: dill<0.3.9,>=0.3.0 in /usr/local/lib/python3.1
Requirement already satisfied: pandas in /usr/local/lib/python3.12/dist-packa
Requirement already satisfied: requests>=2.32.2 in /usr/local/lib/python3.12/
Requirement already satisfied: xxhash in /usr/local/lib/python3.12/dist-packa
Requirement already satisfied: multiprocess<0.70.17 in /usr/local/lib/python3
Requirement already satisfied: fsspec<=2025.3.0,>=2023.1.0 in /usr/local/lib/
Requirement already satisfied: aiohttp!=4.0.0a0,!=4.0.0a1 in /usr/local/lib/p
Requirement already satisfied: hf-xet<2.0.0,>=1.4.2 in /usr/local/lib/python3
Requirement already satisfied: httpx<1,>=0.23.0 in /usr/local/lib/python3.12/
Requirement already satisfied: typing-extensions>=4.1.0 in /usr/local/lib/pyt
Requirement already satisfied: charset_normalizer<4,>=2 in /usr/local/lib/pyt
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.12/dist
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.1
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.1
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/pytho
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/di
Requirement already satisfied: click>=8.2.1 in /usr/local/lib/python3.12/dist
Requirement already satisfied: shellingham>=1.3.0 in /usr/local/lib/python3.1
Requirement already satisfied: rich>=12.3.0 in /usr/local/lib/python3.12/dist
Requirement already satisfied: annotated-doc>=0.0.2 in /usr/local/lib/python3
Requirement already satisfied: aiohappyeyeballs>=2.5.0 in /usr/local/lib/pyth
Requirement already satisfied: aiosignal>=1.4.0 in /usr/local/lib/python3.12/
Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.12/dis
Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.12
Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.
Requirement already satisfied: propcache>=0.2.0 in /usr/local/lib/python3.12/
Requirement already satisfied: yarl<2.0,>=1.17.0 in /usr/local/lib/python3.12
Requirement already satisfied: anyio in /usr/local/lib/python3.12/dist-packag
```

```
Requirement already satisfied: httpcore==1.* in /usr/local/lib/python3.12/dis
Requirement already satisfied: h11>=0.16 in /usr/local/lib/python3.12/dist-pa
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-pac
Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/pyth
Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.12/dist-p
```

```
from transformers import BlipProcessor, BlipForQuestionAnswering

processor = BlipProcessor.from_pretrained("Salesforce/blip-vqa-base")
model = BlipForQuestionAnswering.from_pretrained("Salesforce/blip-vqa-ba
```

```
Loading weights: 100%                                    788/788 [00:00<00:00, 12537.60it/s]
BlipForQuestionAnswering LOAD REPORT from: Salesforce/blip-vqa-base
Key                                        | Status     |  |
-------------------------------------------+------------+--+-
text_encoder.embeddings.position_ids       | UNEXPECTED |  |
text_decoder.bert.embeddings.position_ids  | UNEXPECTED |  |

Notes:
- UNEXPECTED:   can be ignored when loading from different task/architecture;
```

```
data = [
  {
    "image": "/content/sample1.jpg",
    "question": "What is the total amount?",
    "answer": "500"
  }
]
```

```
from torch.utils.data import Dataset
from PIL import Image

class VQADataset(Dataset):
    def __init__(self, data, processor):
        self.data = data
        self.processor = processor

    def __len__(self):
        return len(self.data)

    def __getitem__(self, idx):
        item = self.data[idx]

        image = Image.open(item["image"]).convert("RGB")
        question = item["question"]
        answer = item["answer"]

        encoding = self.processor(image, question, return_tensors="pt")

        labels = self.processor.tokenizer(answer, return_tensors="pt").i

        encoding["labels"] = labels
```

```

        return {k: v.squeeze() for k, v in encoding.items()}
```

```
from torch.utils.data import DataLoader

dataset = VQADataset(data, processor)
loader = DataLoader(dataset, batch_size=2, shuffle=True)
```

```
import torch

optimizer = torch.optim.AdamW(model.parameters(), lr=1e-4)

model.train()
```

```
        )
      )
      (cls): BlipTextOnlyMLMHead(
        (predictions): BlipTextLMPredictionHead(
          (transform): BlipTextPredictionHeadTransform(
            (dense): Linear(in_features=768, out_features=768, bias=True)
            (transform_act_fn): GELUActivation()
            (LayerNorm): LayerNorm((768,), eps=1e-12,
elementwise_affine=True)
          )
          (decoder): Linear(in_features=768, out_features=30524, bias=True)
        )
      )
    )
  )
```

```
model.eval()

from PIL import Image

image = Image.open("/content/sample1.jpg").convert("RGB")
question = "What is the total amount?"

inputs = processor(image, question, return_tensors="pt")
out = model.generate(**inputs)

print(processor.decode(out[0], skip_special_tokens=True))
```

```
---------------------------------------------------------------------------
FileNotFoundError                         Traceback (most recent call last)
/tmp/ipykernel_38447/833523054.py in <cell line: 0>()
      3 from PIL import Image
      4
----> 5 image = Image.open("/content/sample1.jpg").convert("RGB")
      6 question = "What is the total amount?"
      7

/usr/local/lib/python3.12/dist-packages/PIL/Image.py in open(fp, mode,
formats)
   3511     if is_path(fp):
   3512         filename = os.fspath(fp)
-> 3513         fp = builtins.open(filename, "rb")
   3514         exclusive_fp = True
   3515     else:

FileNotFoundError: [Errno 2] No such file or directory:
'/content/sample1.jpg'
```

Next steps: Explain error

```
from google.colab import files
uploaded = files.upload()
```

Choose files Screenshot … 172539.png

Screenshot 2026-03-29 172539.png(image/png) - 57114 bytes, last modified: 29/03/2026 - 100% done

Saving Screenshot 2026-03-29 172539.png to Screenshot 2026-03-29 172539.png

```
Start coding or generate with AI.
```

```
import os
print(os.listdir('/content'))
```

```
['.config', 'Screenshot 2026-03-29 172539.png', 'sample_data']
```

```
image = Image.open("/content/Screenshot 2026-03-29 172539.png").convert(
```

```
Image.open("/content/Screenshot 2026-03-29 172539.png")
```

```
from transformers import BlipProcessor, BlipForQuestionAnswering
from PIL import Image
import requests

# Load processor and model
processor = BlipProcessor.from_pretrained("Salesforce/blip-vqa-base")
model = BlipForQuestionAnswering.from_pretrained("Salesforce/blip-vqa-base")

# Load an image
img_url = "https://storage.googleapis.com/sfr-vision-language-research/BLIP/demo.jpg"
image = Image.open(requests.get(img_url, stream=True).raw).convert("RGB")

# Ask a question
question = "What is in the image?"
inputs = processor(image, question, return_tensors="pt")

# Generate answer
out = model.generate(**inputs)
print(processor.decode(out[0], skip_special_tokens=True))
```

```
model.eval()

from PIL import Image

image = Image.open("/content/Screenshot 2026-03-29 172539.png").convert(
question = "What is the image?"

inputs = processor(image, question, return_tensors="pt")
out = model.generate(**inputs)

print(processor.decode(out[0], skip_special_tokens=True))
```

```
computer code
```

```
from google.colab import files
uploaded = files.upload()

```

Choose files Screenshot … 091259.png

Screenshot 2026-03-30 091259.png(image/png) - 338978 bytes, last modified: 30/03/2026 - 100% done

Saving Screenshot 2026-03-30 091259.png to Screenshot 2026-03-30 091259.png

```
import os
print(os.listdir('/content'))
```

```
['.config', 'Screenshot 2026-03-30 091259.png', 'sample_data']
```

```
from PIL import Image

image = Image.open("/content/Screenshot 2026-03-30 091259.png").convert(
```

```
Image.open("/content/Screenshot 2026-03-30 091259.png")
```

Schedules to the Balance Sheet as at March 31, 2011 and to the Profit and Loss Account for the year ended March 31, 2011

```
model.eval()

from PIL import Image

image = Image.open("/content/Screenshot 2026-03-30 091259.png").convert("
question = "What is the profit percantage?"

inputs = processor(image, question, return_tensors="pt")
out = model.generate(**inputs)

print(processor.decode(out[0], skip_special_tokens=True))
```

$ 5. 00

1.	Net Profit after tax for the year (Rs.)	3,51,19,943	*5,19,15,944*
2.	Number of Equity Shares	31,00,000	*31,00,000*
3.	Basic Earning Per Share (Rs.)	11.33	*16.75*

Start coding or generate with AI.

Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.